

# Attention Based Models for Keyword Spotting

Riccardo Mazziere

**Abstract**—In this work we explore a variety of Deep neural network architectures, with particular focus on the attention mechanism, for the Keyword Spotting task (KWS). In recent years, attention based models have proven to be successful in a wide variety of domains, including image recognition, neural machine translation and speech recognition. We analyze recent attention based architectures: we first focus on the model proposed by de Andrade et al, and explore the role of the size of filters in the convolutional part of the model.

**Index Terms**—Keyword Spotting, Convolutional Neural Networks, Recurrent Neural Networks, Attention Mechanism, Transformers.

## I. INTRODUCTION

The Keyword Spotting (KWS) task consists in the detection of a certain predetermined set of keywords from a stream of user utterances. In recent years, this problem has become increasingly popular and, with the rapid development of mobile devices, it is now playing an important role in human-computer interaction, as well as encouraging the adoption of hands-free interfaces for a wide variety of use cases, that can range from “smart home” devices like Amazon Alexa, to virtual assistants such as Apple’s Siri or Google Assistant. KWS systems are typically required to continuously listen to user inputs on mobile devices, which are highly constrained in their memory and compute capabilities: this restriction has encouraged the development of systems able to achieve highly accurate results but still maintaining a small memory and computational footprint.

At present, Deep Learning (DL) techniques represent the state of the art approach for the KWS problem and have proved to give good results on the tradeoff between model lightness and model performance [1] [2] [3]. In the literature, a lot of DL models have been presented, ranging from simple fully connected feed forward networks [1], to both shallow and very deep Convolutional Neural Networks (CNN) [4] [2] [5] [6]. Especially in recent years, mostly thanks to the work by Vaswani et al. [7], interest towards attention-based architectures has grown dramatically. Indeed, recent works in machine learning have shown that models incorporating attention are able to provide groundbreaking results in a wide variety of domains [8] [9] [10] [11] [12]. Another charming feature of attention-based systems is their inherent interpretability: by visualizing the attention weights, one can directly see to which part of a specific input the model was paying attention when performing inference tasks. Indeed, Explainable AI (XAI) is quickly becoming an hot topic in modern machine learning research, and the development of models capable of giving some sort of explanation for their predictions will in time become more and more desirable, if not required [13].

Motivated by those increasing trends, in this work we first give a review of the existing applications of the attention mechanism for the KWS task; then, we propose different variations of an hybrid convolutional and recurrent model based on attention, which was first introduced by de Andrade et al. [14]; this is done mostly to harness the importance of each block in the original model, and to see if a lighter model in terms of memory footprint can be derived. We summarize the proposed architectures as follows:

- We explore the role of the convolutional part of the model, by both removing it and enhancing it with a more complex architecture based on residual layers;
- We explore the role of the Attention layer by both using more query vectors<sup>1</sup> instead of a single one, and by employing a Multi Head Attention layer [7], with a varying number of heads.

We then compare the obtained performances also in terms of number of parameters for each model. Specifically, we find that the convolutional part can be replaced...

## II. RELATED WORK

### A. Foundations and state of the art

In recent years, machine learning techniques have proven to be the de-facto standard for approaching the KWS problem. Such models typically perform segmentation of the audio sequence on the time domain and extract log-mel scale spectrograms or mel-frequency cepstral coefficients (MFCC) [15] from each frame. Those are then used as input feature vectors for the models. One of the first works exploring deep neural networks for the KWS task is from Chen et al [1]: the authors explore small footprint fully connected architectures and show how they improved performances with respect to the baseline HMM models. This kind of architecture accounted for the sequentiality of audio data by stacking feature vectors of adjacent audio frames: while this gives good results compared to the baseline, it is a very simplistic way to model sequential data. Sainath and Parada [2] approached the problem by using CNNs, a more fitting class of models which are able, by design, to capture several essential features of speech data (like input topology or translational invariance) and at the same time having a much smaller memory footprint due to parameter sharing. Indeed, CNNs have proven to be very successful for KWS: the current state of the art has recently been achieved in [16], where the authors introduce Broadcasted Residual Networks (BC-ResNets), a particular class of ResNets that are able to capture both 1D and 2D

<sup>1</sup>In this work, we use the *query*, *key*, *values* terminology introduced in [7].

features thanks to the introduction of a new kind of residual block.

### B. Models based on Attention

As a premise, it is useful to note that in this applications of attention, the query, key and value vectors are all represented by the feature vectors from the input sequence; therefore when referring to the attention mechanism, we will actually implicitly mean self-attention [17]. Specifically, dot-product attention is adopted [18]. In [14], De Andreade et al. propose an attention based convolutional recurrent neural network (Att-RNN), which takes as input a mel-scale spectrogram and extracts features in the frequency dimension with two convolutional layers. Then, such features are given as input to two bidirectional LSTM [19] layers, in order to extract long range dependencies from the inherently sequential audio data. The last<sup>2</sup> output feature from the LSTM layer is then transformed with a linear dense layer and used as a query vector for the attention mechanism. Finally, the sum of the LSTM outputs, weighted with respect to the attention scores, are given as input to a final dense MLP, which performs the classification. An immediate drawback of this approach is its impossibility to function in a streaming fashion (at least without introducing delay), since bidirectional recurrent layers require the whole sequence of data to work. Furthermore, by using a single vector representation as a query vector representative for the whole sequence, this could act as an information bottleneck. Nevertheless, the introduction of the attention mechanisms allows us to inspect the sections of audio which were responsible for the classification, increasing model transparency. In [3] the authors propose a more modern variant of the same architecture (referred as MHAtt-RNN): LSTM is substituted by GRU [20], and the attention mechanism is replaced by a multi-head attention layer using 4 heads. Despite increasing accuracy with respect to Att-RNN, this kind of models can be very heavy in terms of memory footprint, especially if using a high encoding dimension for the multi-head attention layer. In the case of [3], the number of parameters is 743000 which is more than four times the number of parameters of Att-RNN.

The latest contribution in attention based models for KWS comes from [21], where the authors, inspired by the success of the newly introduced Vision Transformer (ViT) [8] for computer vision tasks, propose an adaptation of such architecture for keyword spotting, called the Keyword Transformer (KWT). The results were suprising: despite ViT proved to be competitive only when supported by pre-training on large datasets, KWT outperformed more complex models based on mixes between CNN, RNN and attention, even when trained on relatively small datasets, like the Google Speech Commands dataset [22]. Even in this case, despite good

performances, it might not be feasible to integrate the KWT model on small mobile devices: indeed, the best performing variation of the KWT that was proposed is relying on a very deep architecture (12 Transformer encoder layers), which counts more than 5 million parameters.

By contrast, we propose lighter attention based architectures, taking Att-RNN as the baseline.

## III. EXPERIMENTAL SETUP

### A. Signals and Features

Each model was trained on the Google Speech Commands dataset V2 [22], which consists on a total of 105829 user utterances of a total of 35 keywords. Each utterance is stored as a one second long<sup>3</sup> WAVE file, sampled at a 16KHz rate. From each signal, 40 Mel Frequency Cepstral Coefficients are computed, using a window size of 25ms (400 samples) and a hop size of 10ms (160 samples). With this approach, the input for the network consists in 3D tensors with one outer channel, where width represents the time domain and height represents the frequency domain. Following the approach from past literature, the models are trained for two different tasks, which are the following:

- **12kws** task: the model must discriminate among 12 different keywords: “yes”, “no”, “up”, “down”, “left”, “right”, “on”, “off”, “stop”, “go”, unknown or silence. The unknown keywords are randomly chosen from the set of remaining keywords and the silence samples consist in one second long crops, randomly extracted from the noise files provided by the dataset<sup>4</sup>. Since the total amount of words belonging to the “unknown” class was much more than number of representatives for each of the other keywords, we decided to randomly extract them to create a balanced dataset. In this way, we have 36921 samples for training, 4443 for validation and 4888 for testing.
- **35kws** task: the model must discriminate among all the 35 keywords present in the dataset. No unknown class or silence class is introduced, therefore, for this task, the entire dataset is used. The training set consists in 84843 samples, the validation set in 9981 samples and the test set in 11005 samples.

Following Google’s suggestions from [22], the dataset is split in 80% for training set, 10% for validation set and 10% for test set, in such a way that the same speakers are never present in two different splits.

For each task, the training set is augmented, following existing approaches from the literature. Specifically, the augmentation process follows this order:

- 1) Each signal is randomly shifted left or right (zero padding the remaining portion), by  $x$  samples, where  $x$  is drawn from a uniform distribution  $U(0, 100)$ ;

<sup>2</sup>In the paper, the authors report to use the middle output, but looking at their implementation, they use the last one. Either way, as they point out, any output of the LSTM layer should work well as a query vector, since the bidirectional LSTM layer should be able to summarize the whole sequence in any of its outputs.

<sup>3</sup>Some clips are a bit shorter than one second: in those cases, the clips are zero padded towards the end.

<sup>4</sup>Those consist in 6 files containing noisy background sounds, both artificially generated and recorded from real environments.

- 2) Samples belonging to the “silence” class do not remain the same across epochs, but are randomly generated each time;
- 3) With a probability of 0.8, each signal is mixed with a randomly generated background noise, which is multiplied by a factor drawn from  $U(0, 0.2)$ ;
- 4) After the conversion in MFCC the features are augmented using SpecAugment [23], a simple data augmentation method which consists in adding randomly sized frequency and time masks to the feature matrix: this is done in order to render the model more robust to partial loss of frequency information or of small sections of speech.

The image resulting from the last step constitutes the actual input for the model. Specifically, each image is a  $98 \times 40$  matrix, where the  $i$ -th row represents the MFCCs of the  $i$ -th time frame.

All the project was carried out using an NVIDIA GeForce GTX 1060 6GB GPU and an Intel i5-6400 CPU, on a Linux machine. All models were built and trained using TensorFlow [24]. All source code is available on Github<sup>5</sup>.

### B. Input Pipeline

Since data augmentation is applied, storing the entire dataset in memory would not be possible: this would result in the same data being reused across epochs. For this reason, a core part of this work was to build an efficient input pipeline that could handle data augmentation on the fly, during training. In this section we explain how the data generation process described in Section III-A works. The augmentation takes place in two different moments during training:

- 1) Phase 1: this is the phase when random noise samples are extracted, and each sample is shifted. This is performed by the CPU, while the GPU is performing training.
- 2) Phase 2: this phase is performed inside the model by the GPU, with a series of preprocessing layers. Specifically, several custom preprocessing layers were built:
  - a) RandomNoiseAugment
  - b) LogMelSpectrogram
  - c) MFCC
  - d) SpecAugment

Note that in Phase 1, each operation is performed one sample at a time, while in Phase 2 each operation is performed in parallel on an entire batch of samples: this results in a much more optimized implementation. Given the hardware setup with which the project was carried on, this framework was a necessity: just computing the MFCC for each sample on CPU would cause a performance bottleneck, wasting the computational capabilities of the GPU. In Figure 1 we report a detailed diagram, showing all the steps in the input pipeline for the generation of the training set.

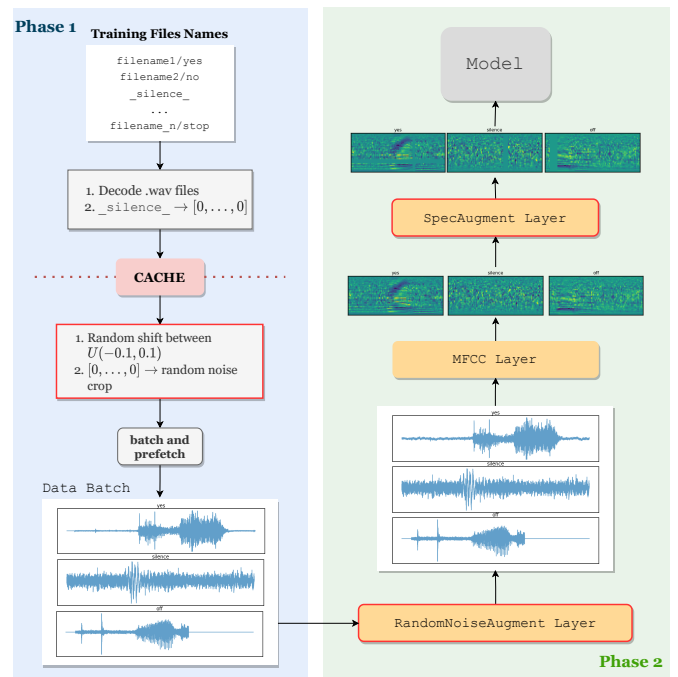


Fig. 1: Detailed description of the system that generates the training data. In the cache block, the data that was produced until that moment is saved to file. Each step which is performed before the caching operation happens only one time, during the first iteration of the dataset; on all the successive iterations, the data is read from the cached file. For the validation and test sets, the augmentation is not applied. Furthermore, to ensure complete separation between training and validation/test sets, the noise samples were generated from different portions of the noise files based on whether they were used for training or for validation or testing.

### C. Learning Framework

Descrivere nel dettaglio i modelli [...]

Each model was trained for 50 epochs, using early stopping and reduce lr on plateau

## IV. RESULTS

In Table 1 we report the different accuracies...

In Figure 2 we plot the test set accuracy with their number of parameters, both for the 12KWS and 35KWS tasks.


### A. Attention Plots

We can plot the attention scores for our models to see which portions of the audio files were catching the attention of the model... In Figure 3 we can make a comparison between multi head and non multi head...


## V. CONCLUDING REMARKS

[illegible]

<sup>5</sup>Link here.



(a) A subfigure



(b) A subfigure

Fig. 2: A figure with two subfigures

TABLE 1: Table with results

Model	Param.	Mult.	Acc.
res8-narrow[7]	20 K	5.65M	90.1%
res15-narrow[7]	43 K	160M	94.0%
res8[7]	111 K	30M	94.1%
res15[7]	239 K	894M	95.8%
DS-CNN-S[5]	24 K	5.4M	94.4%
DS-CNN-M[5]	140 K	19.8M	94.9%
DS-CNN-L[5]	420 K	56.9M	95.4%
TC-ResNet8[8]	66 K	1.12M	96.1%
TC-ResNet8-1.5[8]	145 K	2.20M	96.2%
TC-ResNet14[8]	137 K	2.02M	96.2%
TC-ResNet14-1.5[8]	305 K	4.13M	<b>96.6%</b>
TENet6-narrow	<b>17K</b>	<b>553K</b>	96.0%
TENet12-narrow	31 K	895 K	96.3%
TENet6	54 K	1.68M	96.4%
TENet12	100 K	2.90M	<b>96.6%</b>

## REFERENCES

- [1] G. Chen, C. Parada, and G. Heigold, “Small-footprint keyword spotting using deep neural networks,” in *2014 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 4087–4091, 2014.
- [2] T. Sainath and C. Parada, “Convolutional neural networks for small-footprint keyword spotting,” in *INTERSPEECH*, 2015.
- [3] O. Rybakov, N. Kononenko, N. Subrahmanya, M. Visontai, and S. Laurenzo, “Streaming keyword spotting on mobile devices,” *Interspeech 2020*, Oct 2020.
- [4] R. Tang and J. Lin, “Deep residual learning for small-footprint keyword spotting,” in *ICASSP 2020-2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 5484–5488, IEEE, 2018.
- [5] S. Mittermaier, L. Kürzinger, B. Waschneck, and G. Rigoll, “Small-footprint keyword spotting on raw audio data with sinc-convolutions,” in *ICASSP 2020-2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 7454–7458, IEEE, 2020.

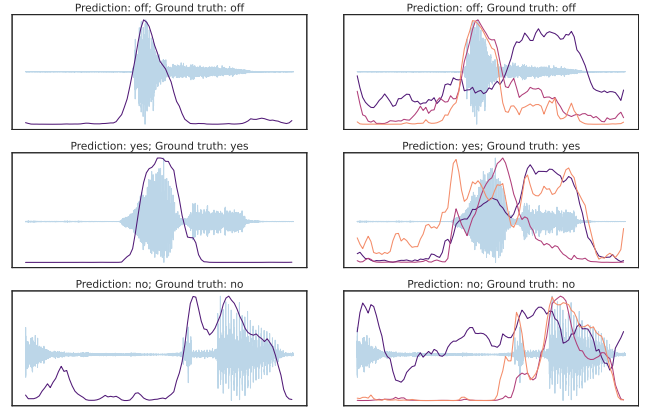


Fig. 3: Comparison between attention scores from Att-RNN model (left) and MHatt-RNN with three heads (right), on the words *off*, *yes* and *no*. We can see that Att-RNN has only one head, so one set of attention scores per prediction is computed. MHatt-RNN instead computes one set of attention weights per head: here we visualize the attention scores for each head. In these examples, we can see how each head learns to pay attention to different phonemes of the same word. In the first example, Att-RNN pays attention only to the first phoneme /o/, while MHatt-RNN has two heads paying attention to /o/ and one paying attention to /f/. In the second example, a similar thing happens: Att-RNN pays attention just at the /je/ while MHatt-RNN has different heads concentrating both on /je/ and /s/. The third example presents a noise at the beginning which is not part of the spoken word: Att-RNN has its attention drawn a bit, while two of three heads from MHatt-RNN learn to completely ignore it.

- [6] S. Choi, S. Seo, B. Shin, H. Byun, M. Kersner, B. Kim, D. Kim, and S. Ha, “Temporal convolution for real-time keyword spotting on mobile devices,” *arXiv preprint arXiv:1904.03814*, 2019.

- [7] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in neural information processing systems*, pp. 5998–6008, 2017.
- [8] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, *et al.*, "An image is worth 16x16 words: Transformers for image recognition at scale," *arXiv preprint arXiv:2010.11929*, 2020.
- [9] H. Touvron, M. Cord, M. Douze, F. Massa, A. Sablayrolles, and H. Jégou, "Training data-efficient image transformers & distillation through attention," in *International Conference on Machine Learning*, pp. 10347–10357, PMLR, 2021.
- [10] A. Gulati, J. Qin, C.-C. Chiu, N. Parmar, Y. Zhang, J. Yu, W. Han, S. Wang, Z. Zhang, Y. Wu, *et al.*, "Conformer: Convolution-augmented transformer for speech recognition," *arXiv preprint arXiv:2005.08100*, 2020.
- [11] M. Kumar, D. Weissenborn, and N. Kalchbrenner, "Colorization transformer," *arXiv preprint arXiv:2102.04432*, 2021.
- [12] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "Bert: Pre-training of deep bidirectional transformers for language understanding," in *NAACL*, 2019.
- [13] B. Goodman and S. Flaxman, "European union regulations on algorithmic decision-making and a 'right to explanation'," *AI Magazine*, vol. 38, pp. 50–57, Oct 2017.
- [14] D. C. de Andrade, S. Leo, M. L. D. S. Viana, and C. Bernkopf, "A neural attention model for speech command recognition," *arXiv preprint arXiv:1808.08929*, 2018.
- [15] S. Davis and P. Mermelstein, "Comparison of parametric representations for monosyllabic word recognition in continuously spoken sentences," *IEEE Transactions on Acoustics, Speech, and Signal Processing*, vol. 28, no. 4, pp. 357–366, 1980.
- [16] B. Kim, S. Chang, J. Lee, and D. Sung, "Broadcasted residual learning for efficient keyword spotting," *arXiv preprint arXiv:2106.04140*, 2021.
- [17] J. Cheng, L. Dong, and M. Lapata, "Long short-term memory-networks for machine reading," *arXiv preprint arXiv:1601.06733*, 2016.
- [18] M.-T. Luong, H. Pham, and C. D. Manning, "Effective approaches to attention-based neural machine translation," *arXiv preprint arXiv:1508.04025*, 2015.
- [19] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [20] K. Cho, B. V. Merriënboer, Çağlar Gülçehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, "Learning phrase representations using rnn encoder-decoder for statistical machine translation," in *EMNLP*, 2014.
- [21] A. Berg, M. O'Connor, and M. T. Cruz, "Keyword transformer: A self-attention model for keyword spotting," *arXiv preprint arXiv:2104.00769*, 2021.
- [22] P. Warden, "Speech commands: A dataset for limited-vocabulary speech recognition," *arXiv preprint arXiv:1804.03209*, 2018.
- [23] D. S. Park, W. Chan, Y. Zhang, C.-C. Chiu, B. Zoph, E. D. Cubuk, and Q. V. Le, "SpecAugment: A simple data augmentation method for automatic speech recognition," *arXiv preprint arXiv:1904.08779*, 2019.
- [24] M. Abadi, P. Barham, J. Chen, Z. Chen, A. Davis, J. Dean, M. Devin, S. Ghemawat, G. Irving, M. Isard, M. Kudlur, J. Levenberg, R. Monga, S. Moore, D. Murray, B. Steiner, P. Tucker, V. Vasudevan, P. Warden, M. Wicke, Y. Yu, and X. Zhang, "Tensorflow: A system for large-scale machine learning," in *OSDI*, 2016.